

# Mastering Prompt Engineering: A Step-by-Step Guide

*Want to land a \$100,000+ tech job in 2025? Start by mastering one of the most high-leverage skills today: prompt engineering.* This guide will walk you through the fundamentals of prompt engineering and the **5 key techniques** you need to know – **no degree or prior experience needed**. By the end, you'll understand how to craft effective prompts that make AI models more **accurate, helpful, and aligned** with your goals.

Prompt engineering is the art and science of writing instructions (“prompts”) that guide an AI model to generate the output you want <sup>1</sup>. It has quickly become a core skill for working with AI, because *how* you ask directly shapes *what* you get <sup>2</sup>. A well-crafted prompt can yield more accurate, relevant, and structured responses, reducing the need for edits or follow-ups <sup>3</sup>. In fact, prompt engineering isn’t just a passing trend – it’s now key to making generative AI systems useful, reliable, and safe <sup>4</sup>. Some prompt engineers are even commanding six-figure salaries as companies recognize the value of this skill <sup>5</sup>.

**How do you “engineer” a good prompt?** The secret is understanding a few proven techniques and applying them step by step. Below, we’ll cover **five essential prompt engineering techniques** and how to master each one:

1. **Role Prompting** – Shaping the AI’s persona or voice
2. **Chain-of-Thought Prompting** – Guiding the AI to reason through problems
3. **Zero-Shot & Few-Shot Prompting** – Using examples (or none) to steer the AI
4. **Prompt Chaining** – Breaking complex tasks into multi-step prompt sequences
5. **Constraint & Format Engineering** – Imposing rules and format on the AI’s output

Throughout this guide, we’ll explain each technique, show why it’s useful, and give you practical steps to apply it. Let’s dive in!

## 1. Role Prompting

Role prompting means **assigning a persona, role, or perspective to the AI** before asking your question. In other words, you tell the model *who it is* or *what role it should play*. This simple step can profoundly influence the **tone, style, and focus** of the AI’s response <sup>6</sup>. For example, compare these prompts:

- *Prompt A:* “Explain the laws of thermodynamics.”
- *Prompt B:* “You are a **science teacher** explaining the laws of thermodynamics to a 12-year-old, using simple language and fun examples.”

Prompt B uses role prompting (“You are a science teacher...”) to set a context. The result? The AI will likely give a **simpler, more accessible explanation** with an educational tone, because we’ve framed it as a teacher addressing a young student <sup>6</sup>. Role prompts can align the model’s voice and behavior with a

specific domain or persona – like a legal advisor, a customer support agent, or a cybersecurity analyst <sup>7</sup>. This is especially useful for tasks requiring a particular **tone or expertise** in the answer <sup>7</sup>.

### How to use Role Prompting effectively:

- **State the Role Clearly:** Begin your prompt by saying “*You are [role]...*”. For instance, “You are a helpful travel assistant,” or “Act as an expert financial advisor.” This immediately frames the context for the model <sup>8</sup>.
- **Impart Relevant Traits:** If needed, add specifics about the role’s behavior or style. E.g., “You are a skeptical analyst – focus on potential risks and controversies in your answers.” This guides the model on *how to behave* within that role <sup>7</sup>.
- **Match the Task to the Role:** Choose a role that fits your query. A medical question might use “You are an experienced doctor,” while a coding question could use “You are a senior software engineer.” The model will try to incorporate domain knowledge and tone appropriate to that persona.
- **Be Consistent:** If you’re in a multi-turn conversation, you can remind or maintain the role as you go. In chat-based interfaces (like the ChatGPT API), you might set the system message to establish the role up front <sup>7</sup>.

**Common Pitfall:** Simply assigning a role isn’t always enough – you should also hint *how the role should influence the answer*. For example, just saying “You are a historian.” might not change the output much. Instead, combine it with instruction: “You are a historian. Provide context from the 18th century when answering.” Avoid leaving the role’s impact implicit <sup>9</sup>.

Role prompting helps **control the voice and perspective** of AI outputs. It’s a powerful way to get the model to adopt the right tone (professional, casual, scientific, friendly, etc.) or to leverage domain-specific knowledge. Practice by experimenting with different roles for the same question to see how the answers change – you’ll be surprised how much the tone and depth shift with a simple persona change <sup>6</sup>.

## 2. Chain-of-Thought Prompting

One of the most powerful techniques for complex problems is **Chain-of-Thought (CoT) prompting** – *guiding the model to think step by step*. Instead of asking for the answer directly, you prompt the AI to **show its reasoning process** or break down the problem into intermediate steps <sup>10</sup>. This approach is like saying: “Don’t jump straight to the answer – **work it out logically first**.”

**Why use chain-of-thought?** For tasks that involve reasoning – math problems, logical puzzles, troubleshooting, or any multi-step analysis – having the model explicitly generate a line of reasoning often leads to more accurate results <sup>11</sup>. By seeing or following a “train of thought,” the AI is less likely to make leaps of logic or mistakes. In fact, a simple prompt like “*Let’s think this through step by step*” can drastically improve problem-solving accuracy on challenging questions <sup>11</sup>.

*Chain-of-Thought prompting leads the model to break down problems. In the example above, a standard prompt yields an incorrect answer, but a chain-of-thought prompt (“Let’s think step by step...”) guides the model to reason through the math, resulting in the correct answer* <sup>10</sup> <sup>11</sup>.

### How to apply Chain-of-Thought prompting:

- **Use trigger phrases:** Simply instruct the model to “think step by step” or say “First, let’s consider...”. For example: *“If a store sells apples for \$2 and oranges for \$3, and I have \$11, how many fruits can I buy? Let’s think step by step.”* This cue encourages the AI to outline the reasoning before giving a final answer <sup>12</sup>.
- **Ask for justification:** You can ask the model to explain its reasoning. E.g., “Explain your reasoning before giving the final answer.” This often produces a stepwise explanation followed by the answer.
- **Provide an example of reasoning (Few-shot CoT):** In a few-shot prompt (more on few-shot later), you can show an example where a question is answered with a reasoning process (e.g., a Q&A pair where the solution is worked out in steps). The model will mimic that format for new questions <sup>10</sup>.
- **Use structured reasoning format:** For advanced control, you might use formatting like: “<thinking> ... </thinking><answer> ... </answer>” to explicitly separate reasoning from the final answer <sup>13</sup>. This can help the model keep track of which part is explanation versus conclusion (and some models like Claude respond well to such tags <sup>13</sup>).

**When to use it:** Chain-of-thought prompting is ideal for **any problem that isn’t trivial and benefits from analysis**, such as: multi-step math calculations, logical inference, coding/debugging steps, complex decision making, or scenarios where correctness is critical <sup>14</sup>. It essentially forces the AI to *show its work*. Even if you only need the final answer, guiding the reasoning can reduce errors. As a bonus, you (or your users) can see the thought process, making the result more transparent and easier to double-check.

**Tip:** Chain-of-thought can be combined with other techniques. For example, you might use a **Role + Chain-of-Thought combo**: *“You are a cybersecurity analyst. Below is an incident report. Think step by step through the possible causes and then give your resolution.”* This merges a domain-specific role with logical reasoning guidance – a powerful approach for robust performance <sup>15</sup>.

In summary, **don’t hesitate to ask the AI to think out loud**. It might feel unusual at first (“Why would I tell the computer to *think*?”), but this technique leverages the model’s ability to follow procedural text. With practice, chain-of-thought prompting will become your go-to for tackling tough problems that require more than a one-line answer.

## 3. Zero-Shot and Few-Shot Prompting

Next, let’s talk about how to use **examples (or the lack thereof)** in your prompts. These techniques are about *how much context or demonstration you give the model* when asking a question. We have two ends of a spectrum:

- **Zero-Shot Prompting:** *No examples provided.* You just give an instruction or question, and the model must respond from its own knowledge. For instance: *“Translate the following sentence to French: ‘Hello, how are you?’”* – a direct instruction <sup>16</sup>.
- **Few-Shot Prompting:** *One or more examples are included* in the prompt to show the model what you expect. For example, you might provide a couple of English-to-French sentence translations first, then ask it to translate a new sentence. This “teaches” the model the pattern or format before it answers <sup>17</sup>.

Few-shot prompting leverages a phenomenon called **in-context learning** – the model learns from the examples in the prompt, without any parameter updates. Essentially, you're **demonstrating the task** within the prompt itself. This can dramatically improve performance on tasks that have specific formats or are a bit complex, by giving the model a reference point <sup>10</sup> <sup>16</sup>.

#### When to use Zero-Shot vs Few-Shot:

- **Zero-Shot** is best for **simple or well-known tasks** where the model can handle it with general knowledge <sup>18</sup>. It's quick and doesn't require you to come up with examples. For instance, summarization, straightforward Q&A, or fact recall are often fine zero-shot. *Always start with a clear, unambiguous instruction*, since that's all the model has to go on <sup>19</sup>. You might be surprised – modern models can do a lot zero-shot if the prompt is phrased well.
- **Few-Shot** is helpful when the task is **complex, format-sensitive, or unusual** <sup>18</sup>. By giving 2-5 examples, you provide a pattern for the model to follow <sup>20</sup>. This is great for things like: outputting in a specific format (bullet lists, JSON, etc.), adopting a particular style or tone, or handling tasks with subtle nuances. Few-shot prompts shine in classification, transformation, or creative tasks where the *style* of the output is important. They also help if the model tends to make mistakes zero-shot – an example can clarify what you want.

Consider an example of few-shot prompting for sentiment analysis. A few-shot prompt could be:

```
**Task:** Determine if the sentiment of a sentence is Positive, Negative, or Neutral.  
  
**Examples:**  
"I love the new design of your website!" → Positive  
"This product is terrible and broke on day one." → Negative  
"I think the movie was okay, not great but not bad." → Neutral  
  
**Now classify:**  
"The service at the restaurant was slow and unprofessional." →
```

In this prompt, we gave the instruction (the task) and three example sentences with their sentiments. The model will use these to infer that the format is "*sentence → Sentiment*" and will likely respond with "Negative" for the new sentence. Without those examples (zero-shot), the model might still do fine if asked directly, but the examples **reduce ambiguity and guide the output format**.

#### Tips for effective Zero/Few-Shot Prompting:

- **Clarity is king (Zero-Shot):** When no examples are given, make your instruction as explicit as possible. Specify the task, the perspective, and even the desired format if applicable <sup>21</sup> <sup>22</sup>. For example: *"You are a sentiment analysis model. Classify each sentence as 'Positive', 'Negative', or 'Neutral'. Respond with one word for each."* – this zero-shot prompt sets a role, the task, and output format clearly, which improves the chances of a correct response <sup>23</sup> <sup>22</sup>.
- **Use One-Shot or Few-Shot for format/tone:** If you need a particular output style, show one example (one-shot) or a few. Models excel at imitating the style or format of given examples <sup>24</sup> <sup>25</sup>.

For instance, to get a poem, provide one sample poem; to get an answer in JSON, show an example question and a JSON-formatted answer.

- **Keep examples concise and relevant:** Don't overload the prompt with lengthy or unnecessary examples. Each example should demonstrate *exactly* the behavior you want. Inconsistent or overly complex examples can confuse the model <sup>26</sup>.
- **Separate examples from the query:** Clearly delineate your examples from the actual question/prompt (you can use headings like "Examples:" and "Now X:" as shown above, or special delimiters). This prevents the model from blending the examples with the real task <sup>27</sup> <sup>28</sup>.
- **Test different numbers of shots:** Sometimes one or two examples suffice; other times adding more improves reliability. There's a point of diminishing returns (and remember, long prompts use more tokens), so experiment to see the minimal number of examples needed to get good results.

In summary, **zero-shot prompting** is quick and relies on the model's internal knowledge – use it for straightforward tasks or when you're prototyping. **Few-shot prompting** gives the model a hand by showing it what you expect – use it for trickier tasks, custom formats, or when zero-shot isn't cutting it. Mastering both approaches (and knowing when to use which) will make you a much more effective prompt engineer <sup>18</sup>.

## 4. Prompt Chaining

Prompt chaining is an advanced technique where you **break a complex task into a sequence of simpler prompts**, chaining the outputs together. Instead of trying to get a single AI response to do everything, you structure a **step-by-step workflow**: the result of Prompt 1 feeds into Prompt 2, and so on <sup>29</sup> <sup>30</sup>. This is incredibly useful when a task is too complex to solve in one go or when a single prompt would become too unwieldy.

Think of it like an assembly line: each prompt handles one subtask, and together they build the final answer. For example, suppose you have a huge document and you need an AI to answer a question about it. A prompt chain approach might be: 1. **Prompt 1:** "Read the following document and extract the key facts relevant to [the question]." → *AI outputs a list of relevant facts or quotes* <sup>31</sup> <sup>32</sup>. 2. **Prompt 2:** "Using the extracted facts, answer the question in a concise paragraph." → *AI uses the facts from Prompt 1's output to produce a focused answer* <sup>33</sup>.

By chaining prompts, the model first concentrates on *finding information*, then on *formulating an answer*. Each step is simpler and more focused than doing everything at once. This often **improves accuracy, clarity, and reliability** of the final output <sup>34</sup> <sup>35</sup>. It also makes the process more transparent – you can inspect and tweak the intermediate outputs if something goes wrong, rather than having a single opaque answer.

**When to use Prompt Chaining:** Use it for **multi-step problems or complex workflows** where a one-shot prompt might fail or produce errors <sup>36</sup>. Some scenarios: - **Research Synthesis:** e.g., Prompt 1 to gather data or sources, Prompt 2 to summarize or analyze them <sup>37</sup>. - **Long Documents/Q&A:** e.g., first prompt to find relevant sections, second prompt to answer using those sections <sup>31</sup> <sup>38</sup>. - **Creative Content Generation:** e.g., Prompt 1 to brainstorm ideas, Prompt 2 to outline, Prompt 3 to draft, Prompt 4 to edit – a chain ensuring each stage is handled thoroughly <sup>39</sup> <sup>40</sup>. - **Data Transformation Pipelines:** e.g., break a task like "translate, extract info, format output" into separate prompts for each step <sup>41</sup> <sup>42</sup>. This reduces the cognitive load on the model at each step.

### How to build a prompt chain (step-by-step):

1. **Identify the Subtasks:** Break down the overall goal into smaller, logical subtasks <sup>43</sup> <sup>42</sup>. Ask: what needs to happen first, second, third? Make each sub-goal as clear and simple as possible.
2. **Design Each Prompt:** For each subtask, write a prompt that accomplishes just that piece. Be explicit about the expected output of that step. For example, "List the names of all characters in the story" (to extract info) or "Summarize the above text in 3 bullet points" (to transform info).
3. **Ensure Clarity and Context:** When moving to the next prompt in the chain, provide the necessary context. This might mean feeding the AI the output from the previous step (you can literally copy the text, or in a programmatic setting, pass it programmatically). Start each prompt with any needed instructions plus the content from prior steps (if needed) clearly separated (you can use delimiters like ##### or quotes to indicate the AI's prior output) <sup>38</sup> <sup>44</sup>.
4. **Single-Task Focus:** Keep each prompt focused on *one task only*. If a prompt is doing two or three things, consider splitting it further. Each "link" in the chain should be straightforward <sup>45</sup>.
5. **Test the Chain:** Run through the prompts in order and see how it performs. You might find one step doesn't give exactly the format you need for the next. Tweak that prompt, or adjust the next prompt to handle the input better.
6. **Iterate and Refine:** Prompt chaining often requires a few iterations. Maybe you need to add an extra step to clean or reformat data, or simplify a prompt for better output. Refining each step will improve the overall result <sup>46</sup> <sup>47</sup>.

One great advantage of prompt chaining is **controllability**. You can insert checks or transformations between steps, ensuring each part is correct before moving on. It also helps with **debugging** your prompt logic – if the final output is off, you can inspect which step in the chain produced a hiccup and fix that, rather than guessing at what went wrong in a giant single prompt <sup>35</sup>.

**Example:** Suppose you're building a prompt chain to analyze a legal contract: - Prompt 1: "Read the contract and extract any clauses related to data privacy." (Output: a list of privacy clauses)

- Prompt 2: "For each clause extracted, explain in plain English what it means and any potential risks." (Output: explanations for each clause)
- Prompt 3: "Summarize the key privacy risks identified in 2-3 bullet points." (Output: a concise bullet list of risks)

By dividing the task, the AI can focus on one layer of the problem at a time – making it more likely to capture details and present them clearly <sup>48</sup> <sup>43</sup>. Meanwhile, as the user or developer, you can review the list of clauses from Prompt 1 or the explanations from Prompt 2 and correct any issues before proceeding.

Prompt chaining is a bit more involved than single prompts, but it pays off with **higher quality and more personalized outputs** <sup>49</sup> <sup>37</sup>. It's a technique that becomes increasingly important as you tackle real-world projects with AI, where tasks often have multiple parts. Start practicing by taking a task you would normally do in one prompt, and split it into two or three. You'll gain a deeper understanding of the model's capabilities at each step and create solutions that are easier to manage and more **robust** in production.

## 5. Constraint & Format Engineering

The final technique is all about **telling the model exactly how you want the answer to be delivered**. Constraint and format engineering means adding instructions that **limit or guide the output's format**,

**style, or content.** Instead of just asking a question, you also specify the answer's shape: whether that's a word limit, a bullet list, a JSON structure, a certain tone, etc. <sup>50</sup>.

By imposing formatting constraints, you "tame" the AI's output to fit your needs <sup>51</sup> <sup>52</sup>. This is incredibly useful for making sure the answer is **structured, consistent, and ready to use**. For example: - "*Explain blockchain in exactly 3 sentences, in a friendly tone.*" – constrains length and style. - "*List 5 reasons in bullet point format why remote work can increase productivity.*" – forces a bullet list format <sup>50</sup>. - "*Respond only with a JSON object containing keys status and solution.*" – forces the output to be JSON, which is great if you plan to feed it into a program <sup>53</sup>.

Why bother with constraints? Because LLMs, left to their own devices, can be verbose, inconsistent, or include unnecessary commentary <sup>54</sup>. By **defining the format or limits**, you steer the model to produce output that is **immediately useful** and easier to parse. In a way, you're making the AI do post-processing for you. For instance, if you need an answer in a table form, it's much easier to tell the AI to output a table than to manually organize the text yourself after. Moreover, structured outputs (JSON, XML, etc.) allow for programmatic processing – crucial if this is part of an application pipeline <sup>55</sup>.

### Techniques for Constraint & Format Engineering:

- **Specify Output Format in the Prompt:** Don't be shy about literally saying "Respond in JSON" or "Answer in 2-3 sentences" or "Provide a numbered list". Models will usually follow that instruction strictly <sup>53</sup> <sup>56</sup>. For example: "*Give the answer as a JSON: { "summary": "...", "rating": "..."} and nothing else.*" By doing this, you can often guarantee that certain fields or structure appear in the response <sup>57</sup> <sup>58</sup>.
- **Set Length or Scope Limits:** If you need brevity or a specific length, include that. "*In under 50 words...*", "*no more than one paragraph,*" or "*exactly 5 bullet points.*" This helps avoid overly long answers and keeps content focused <sup>59</sup>.
- **Constrain Style or Content:** You can instruct the model to avoid certain things or stick to a style. E.g., "Use simple language, suitable for a 5th grader," or "Do not include any apologies or filler text." If you need the model to stay within certain boundaries (like avoiding opinions, or only using metric units), spell that out.
- **Use Structured Prompt Formats:** Some prompt engineers use formats like wrapping the desired answer schema in quotes or code blocks within the prompt so the model knows to fill it in. For instance: "*Fill the following JSON template with appropriate values: { "City": "<City Name>", "Weather": "<Description>" }*". This combines few-shot (providing a template) with format constraint.
- **Reinforce with Role or System Message:** If using a chat system with a system message, you can put formatting rules there as a persistent instruction (e.g., "The assistant should always respond in markdown format with code snippets for code answers.").

**Benefits:** By constraining format, you achieve **consistent and parseable outputs** <sup>60</sup>. This is great for building applications – for example, you can ask for an answer in YAML or JSON and then easily use that output in your code. It also reduces the chance of the model going off on tangents. Studies and practice have shown that adding structure requirements can even improve the relevance and accuracy of responses <sup>61</sup>. Think of it like giving the model a container to fill – it's more likely to put the right stuff in the box and not spill over.

**Pitfalls:** Sometimes, if you over-constrain, the model might say it cannot do that, or it may produce an empty result, etc. Also, adding too many format instructions can increase prompt length (using up tokens). There's a balance to strike between strictness and flexibility. A common trick if the model keeps adding extraneous text is to prepend something like “**IMPORTANT:** Respond *only* with the following format, no explanations.” <sup>62</sup>. Models like GPT-4 or Claude are usually good at following format when told, but lesser models might need more nudging or may still make minor mistakes (like an extra sentence). In critical applications, you might still need to validate the output (for example, if you required JSON, ensure it's valid JSON after getting the response).

In practice, **format engineering** is your way to get an answer that doesn't just answer the question, but does so in a way that's immediately **useful for you**. Whether that's a well-structured report, a snippet of code, a table of results, or a Twitter-length summary, you can usually get the AI to comply by explicitly describing the desired format <sup>50</sup> <sup>60</sup>. It's often as easy as adding one sentence to your prompt, and it can save you a ton of time cleaning up outputs later.

---

## Conclusion & Next Steps

Mastering these five techniques – **Role Prompting, Chain-of-Thought, Zero/Few-Shot, Prompt Chaining, and Format Constraints** – will elevate your prompt engineering skills to a professional level. You've learned how to set the stage for the AI with roles, coax out step-by-step reasoning, provide examples for guidance, split hard problems into manageable prompts, and demand answers in the exact format you need. These are the core tools that expert prompt engineers use daily to get the most out of AI models.

**Practice is key:** Try mixing and matching techniques. Many real-world prompts combine several of these methods for best results <sup>63</sup> <sup>64</sup>. For instance, a complex project prompt might use a Role, provide Few-Shot examples, and request a formatted answer *all at once*. As you experiment, you'll develop an intuition for what the model needs to produce the ideal output. Don't be afraid to iterate – even seasoned practitioners refine their prompts multiple times <sup>65</sup>.

Finally, keep learning and stay curious. The AI field is evolving fast, and prompt engineering strategies continue to improve. What remains constant is the principle that **clear, well-structured prompts yield the best AI responses** <sup>66</sup> <sup>67</sup>. With the step-by-step approach outlined in this guide, you have a roadmap to continually sharpen your prompting skills.

**Now it's your turn** – go ahead and apply these techniques in your next AI project or even in your daily interactions with tools like ChatGPT. You'll not only get better results, but you'll also be building a skill set that's in high demand in today's tech landscape. Happy prompting!

**Sources:** Each technique and tip above is backed by insights from AI research and expert practitioners. For further reading and examples, check out the references linked throughout this guide, such as the *Prompt Engineering Guide* <sup>10</sup> <sup>30</sup>, industry articles <sup>6</sup> <sup>37</sup>, and tutorials <sup>21</sup> <sup>18</sup>. These will deepen your understanding and keep you on the cutting edge of prompt engineering as you progress in your learning journey. Good luck!

---

1 2 3 6 11 12 16 17 24 25 50 56 65 66 67 **Prompt Engineering: The Secret Language Behind Smarter AI** | by Szőcs Zsombor | Medium

<https://medium.com/@szocs.zsombor86/prompt-engineering-the-secret-language-behind-smarter-ai-a08f988edd87>

4 7 8 9 13 14 15 18 26 53 54 59 60 62 63 64 **The Ultimate Guide to Prompt Engineering in 2025** | Lakera – Protecting AI teams that disrupt the world.

<https://www.lakera.ai/blog/prompt-engineering-guide>

5 **Top 10 High-Income AI Skills to Learn in 2025 for Career Growth**

<https://startuptalky.com/top-high-income-ai-skills/>

10 **Chain-of-Thought Prompting** | Prompt Engineering Guide

<https://www.promptingguide.ai/techniques/cot>

19 20 21 22 23 27 28 **Prompt Engineering 101: Understanding Zero-Shot, One-Shot, and Few-Shot** | Codecademy

<https://www.codecademy.com/article/prompt-engineering-101-understanding-zero-shot-one-shot-and-few-shot>

29 41 42 49 **What is prompt chaining?** | IBM

<https://www.ibm.com/think/topics/prompt-chaining>

30 31 32 33 35 38 44 **Prompt Chaining** | Prompt Engineering Guide

[https://www.promptingguide.ai/techniques/prompt\\_chaining](https://www.promptingguide.ai/techniques/prompt_chaining)

34 36 37 39 40 43 45 46 47 48 **Mastering Complex Tasks: The Art of Prompt Chaining** | by NAITIVE | Medium

<https://medium.com/@NAITIVE/mastering-complex-tasks-the-art-of-prompt-chaining-97edc4594757>

51 52 55 57 58 61 **Taming LLM Outputs. Your Guide to Structured Text...** | by Vivien Tran Thien | data from the trenches | Medium

<https://medium.com/data-from-the-trenches/taming-lm-outputs-59a58ee3246d>